

SALTAR Project Overview

Russell Bentley

Stony Brook

2024



Stony Brook University

Table of Contents

1 Introduction

2 DSL

- Stencil System
- Domain
- Boundary Conditions

3 Compilation Techniques

- FFT
- Gaussian Approximation
- Polyhedral Compiling
- Misc Techniques

4 Autotuner

5 Related



Deliverables

- SALTAR
 - DSL
 - Compiler
 - Autotuner
- New approximation algorithms
- Benchmark Suite



Prior Work

- Pachoir compiler [17]
- PLUTO [7], [6]
- Foust compiler (FFT based) [4]
- Halide (C++ DSL) [16]
- Devito (FD focus, useful resource) [14]



Table of Contents

1 Introduction

2 DSL

- Stencil System
- Domain
- Boundary Conditions

3 Compilation Techniques

- FFT
- Guassian Approximation
- Polyhedral Compiling
- Misc Techniques

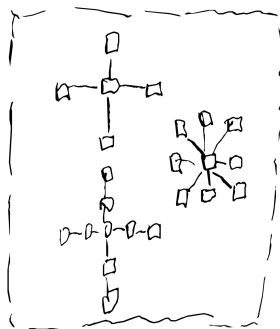
4 Autotuner

5 Related



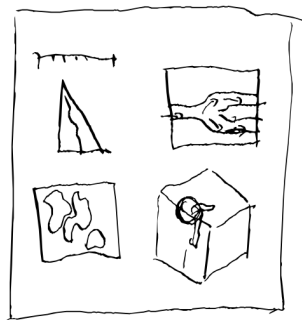
Stencil System

- 1 or more stencil operation
- Each stencil is fully programmable
 - Like a shader, SPMD
 - Limited write access
- Can vary
 - Spatially
 - Temporally
 - By domain composition



Domain

- Node graph / connectivity
- May be geometry driven
- Should be programmable (?)
 - Condensation problems
- Lots of axis aligned boxes in practice
- Might include “coloring” to map different stencils
- Where do want to sample?
- Duration of runtime
 - How often do we sample?



Boundary Conditions

- Fully programmable
 - Highly dependent on domain and stencil system
- May require sampling
 - Outflow boundary conditions
 - Coupling to other simulations

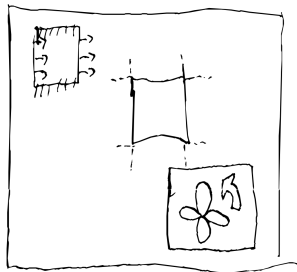


Table of Contents

- 1 Introduction
- 2 DSL
 - Stencil System
 - Domain
 - Boundary Conditions
- 3 Compilation Techniques
 - FFT
 - Gaussian Approximation
 - Polyhedral Compiling
 - Misc Techniques
- 4 Autotuner
- 5 Related



FFT

- FFT algorithms [3] [1] [1]



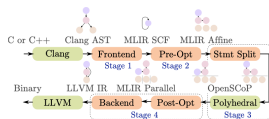
Gaussian Approximation

- Gaussian Approximation [2]



Polyhedral Compilers

- PLUTO [7], [6]
 - C source to source compiler
 - Access optimized with tiling
- LLVM Polygeist [15]
 - C++ / C interface to MLIR [13]
 - MLIR polyhedral optimization passes
- LLVM Polly [10]
 - Polyhedral optimization of LLVM IR



Polygeist workflow



Misc Techniques

- Branch Removal
 - If statements may have arithmetic based alternative
- Vectorizing
- Data layout / Access patterns (AoS \rightarrow SoA)



Table of Contents

- 1 Introduction
- 2 DSL
 - Stencil System
 - Domain
 - Boundary Conditions
- 3 Compilation Techniques
 - FFT
 - Guassian Approximation
 - Polyhedral Compiling
 - Misc Techniques
- 4 Autotuner
- 5 Related



- Profiling?
- Plan Composition
- Algorithm Tradeoffs
- Other hyper parameters?



Table of Contents

- 1 Introduction
- 2 DSL
 - Stencil System
 - Domain
 - Boundary Conditions
- 3 Compilation Techniques
 - FFT
 - Gaussian Approximation
 - Polyhedral Compiling
 - Misc Techniques
- 4 Autotuner
- 5 Related



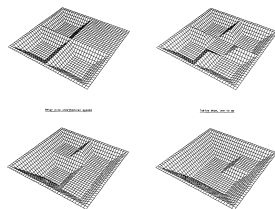
Related Projects

- FFTW [8]
 - Fast fourier transform compiler
 - A dependency(?) for SALTAR
 - Similiar architectural concerns
- Taichi [12]
 - JIT compiler parallel numerical code
 - Optimizes computation over sparse data
 - DSL is based on python.
- Eigen [11]
 - Runtime vs compile time configuration
- FEnics [5]
 - Compiler framework for FEA



Related Algorithms

- HashLife [9]
 - Memoized Algorithm for cellular automata
 - Sensitive to entropy
- Quicklife (Open Source with Golly)
 - Tree based evaluation
 - No hashing



References I

- [1] Zafar Ahmad et al. “A fast algorithm for aperiodic linear stencil computation using Fast Fourier Transforms”. In: *ACM Transactions on Parallel Computing* 10.4 (2023), pp. 1–34.
- [2] Zafar Ahmad et al. “Brief announcement: Faster stencil computations using gaussian approximations”. In: *Proceedings of the 34th ACM Symposium on Parallelism in Algorithms and Architectures*. 2022, pp. 291–293.
- [3] Zafar Ahmad et al. “Fast stencil computations using fast Fourier transforms”. In: *Proceedings of the 33rd ACM Symposium on Parallelism in Algorithms and Architectures*. 2021, pp. 8–21.
- [4] Zafar Ahmad et al. “Fourst: A code generator for FFT-based fast stencil computations”. In: *2022 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. IEEE. 2022, pp. 99–108.



References II

- [5] Igor A Barrata et al. “DOLFINx: The next generation FEniCS problem solving environment”. In: (2023).
- [6] Uday Bondhugula et al. “A practical automatic polyhedral parallelizer and locality optimizer”. In: *Proceedings of the 29th ACM SIGPLAN Conference on Programming Language Design and Implementation*. 2008, pp. 101–113.
- [7] Uday Bondhugula et al. “Automatic transformations for communication-minimized parallelization and locality optimization in the polyhedral model”. In: *Compiler Construction: 17th International Conference, CC 2008, Held as Part of the Joint European Conferences on Theory and Practice of Software, ETAPS 2008, Budapest, Hungary, March 29–April 6, 2008. Proceedings 17*. Springer. 2008, pp. 132–146.



References III

- [8] Matteo Frigo and Steven G Johnson. “The design and implementation of FFTW3”. In: *Proceedings of the IEEE* 93.2 (2005), pp. 216–231.
- [9] R Wm Gosper. “Exploiting regularities in large cellular spaces”. In: *Physica D: Nonlinear Phenomena* 10.1-2 (1984), pp. 75–80.
- [10] Tobias Grosser, Armin Groesslinger, and Christian Lengauer. “Polly—performing polyhedral optimizations on a low-level intermediate representation”. In: *Parallel Processing Letters* 22.04 (2012), p. 1250010.
- [11] Gaël Guennebaud, Benoît Jacob, et al. *Eigen* v3. <http://eigen.tuxfamily.org>. 2010.



References IV

- [12] Yuanming Hu et al. “Taichi: a language for high-performance computation on spatially sparse data structures”. In: *ACM Transactions on Graphics (TOG)* 38.6 (2019), p. 201.
- [13] Chris Lattner et al. “MLIR: Scaling compiler infrastructure for domain specific computation”. In: *2021 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*. IEEE. 2021, pp. 2–14.
- [14] Fabio Luporini et al. “Architecture and performance of Devito, a system for automated stencil computation”. In: *ACM Transactions on Mathematical Software (TOMS)* 46.1 (2020), pp. 1–28.
- [15] William S Moses et al. “Polygeist: Raising C to polyhedral MLIR”. In: *2021 30th International Conference on Parallel Architectures and Compilation Techniques (PACT)*. IEEE. 2021, pp. 45–59.



References V

- [16] Jonathan Ragan-Kelley et al. “Halide: a language and compiler for optimizing parallelism, locality, and recomputation in image processing pipelines”. In: *Acm Sigplan Notices* 48.6 (2013), pp. 519–530.
- [17] Yuan Tang et al. “The Pochoir Stencil Compiler”. In: *Proceedings of the twenty-third annual ACM symposium on Parallelism in algorithms and architectures*. 2011, pp. 117–128.

