MKL Performance Over Time

Russell Bentley

Stony Brook

2025

Stony Brook University

Table of Contents

Sourcing MKL Versions

- Intel only releases from the last year on their software center.
- Releases through 2023 through Intel's apt repository
- Some 2020 releases on Archive.org
- Official MKL python wheels on pypi include dynamic libraries as far back as 2018
  - All benchmarks used dynamic linking as fewer static libraries could not be sourced

Experimental Setup

Benchmarks

- One warm-up iteration
- 16 timed iterations
- Plots show average time
- Error bars show standard deviation
- Sequential and OpenMP variants
  - 80 threads for OpenMP builds
  - 2025.0 build of Intel's OpenMP runtime used for all benchmarks

Machine

- IceLake node on TACC
- Intel Xeon Platinum 8380 (Q2'21)
- 80 cores on two sockets

Table of Contents

dgemm

- Double precision general matrix multiply
- $C = \alpha A \cdot B + \beta C$
- Less noise in sequential runtimes - expected

dgemm - OpenMP

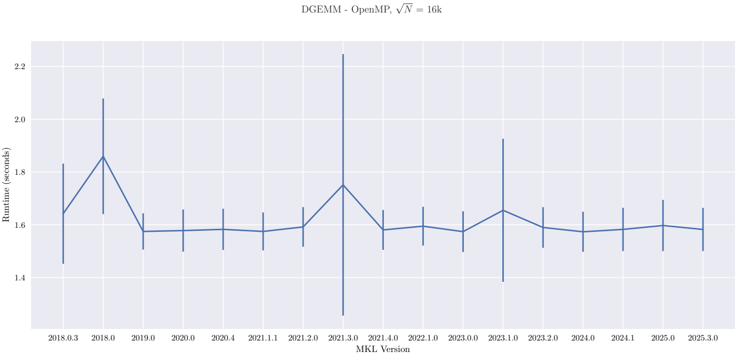

dgemm - Sequential

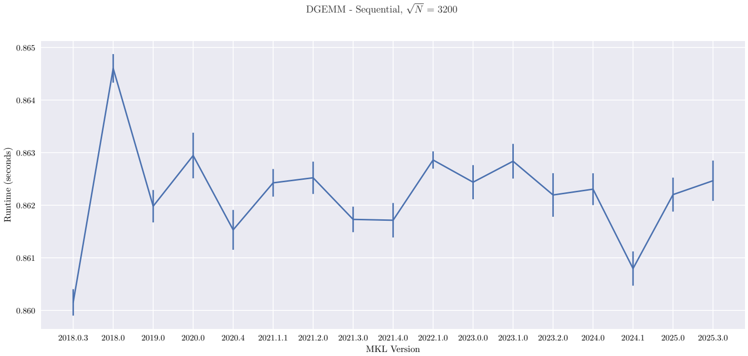

Table of Contents

dgeqrf

- Double precision QR factorization

dgeqrf - OpenMP

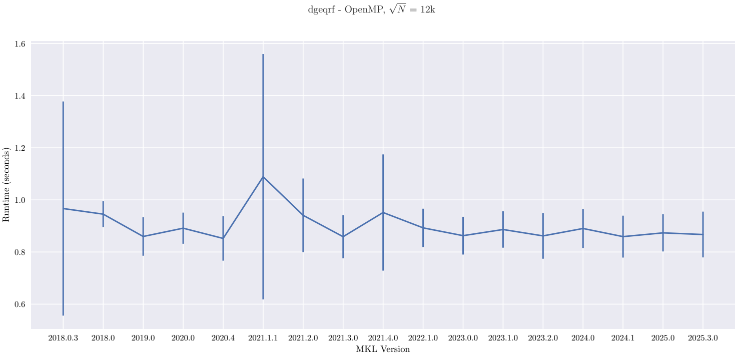

dgeqrf - Sequential

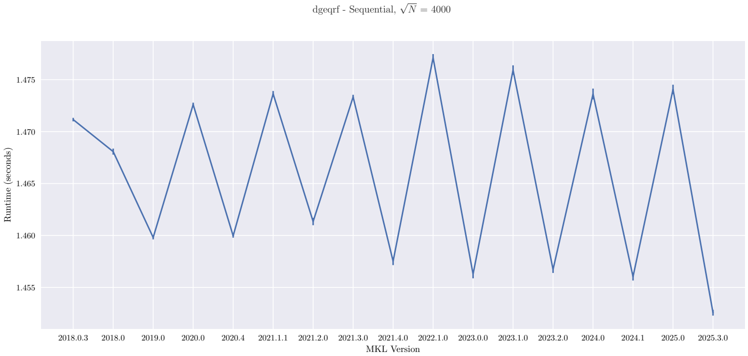

Table of Contents

dtrsm

- Double precision solve for

$$\mathrm{op}(A) \cdot X = \alpha \cdot B$$

dtrsm - OpenMP

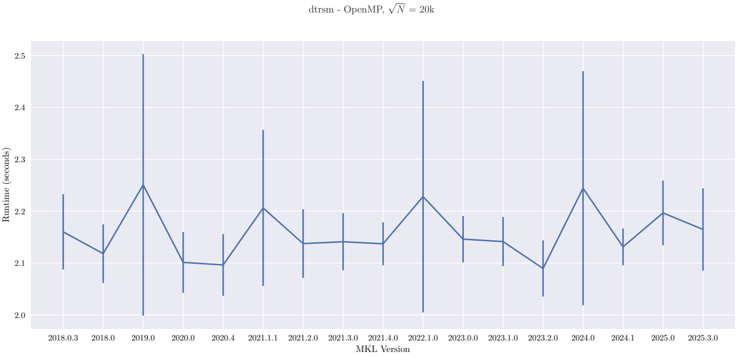

dtrsm - Sequential

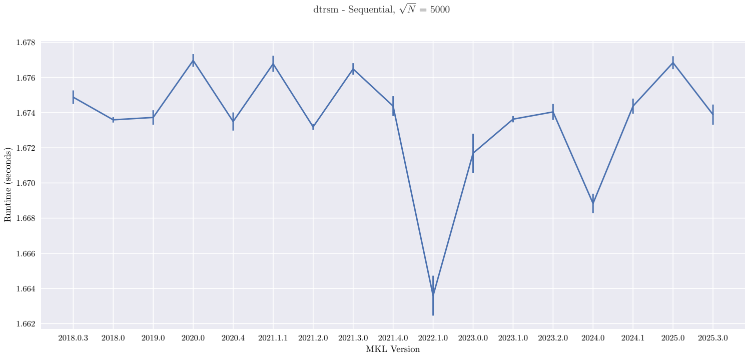

Table of Contents

DFT 2d

- Double precision discrete Fourier transform
  - General improvement over time
- Trend is very clear for sequential
- From late 2018 release notes (build not tested)
  - Improved performance for batched real-to-complex 3D for Intel® Xeon® Processor supporting Intel® Advanced Vector Extensions 512 (Intel® AVX-512) (codename Skylake Server)
  - Improved performance with and without scaling factor across all domains.

DFT 2D - OpenMP

DFT - OpenMP, $\sqrt{N} = 32k$

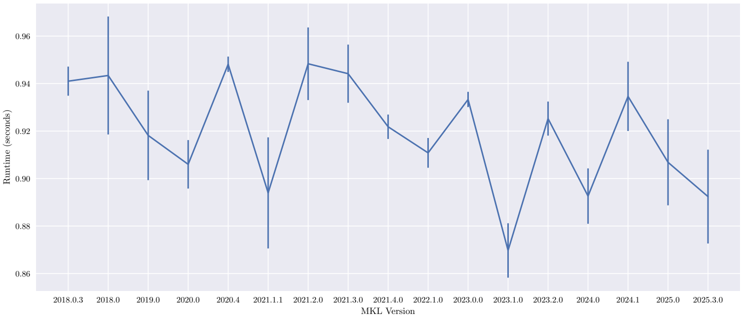

DFT 2D - Sequential

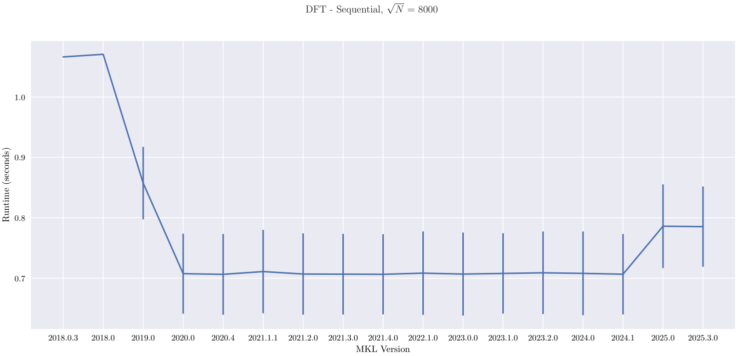

Table of Contents

DFT 2d

- Double precision direct sparse linear solver
- Solve `s3dkt3m2.mtx`
- Improvement starting in 2025
- From release notes:
  - "Improved iterative refinement feature of oneMKL PARDISO"

Pardiso - OpenMP

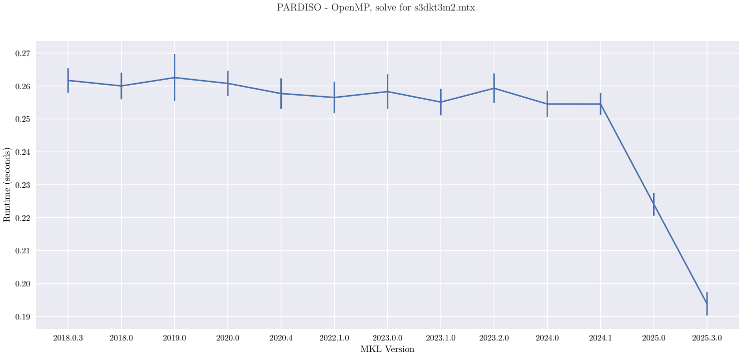

Pardiso - Sequential

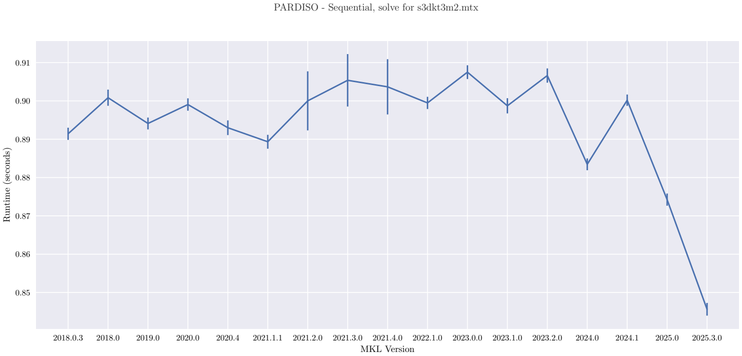

Table of Contents

vsExp

- Single precision element-wise exponential
- $a_i = e^{b_i}$
- Performance has degraded over time

vsExp - OpenMP

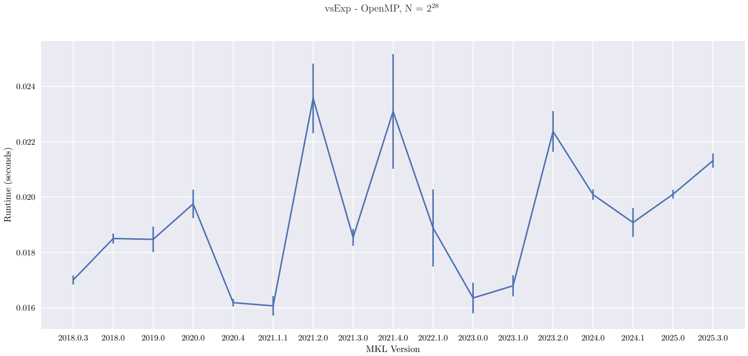

vsExp - Sequential

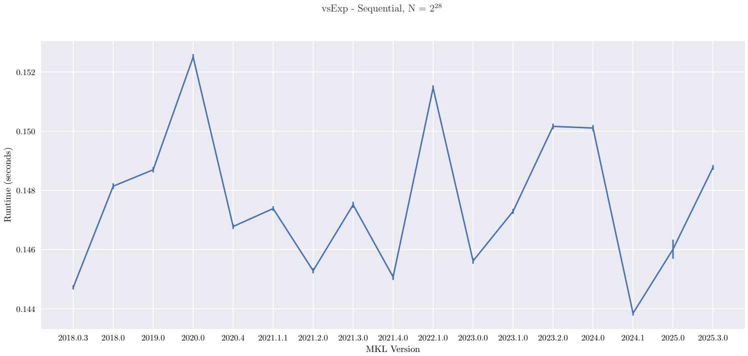

Table of Contents

vdExp

- Double precision element-wise exponential
- $a_i = e^{b_i}$

vdExp - OpenMP

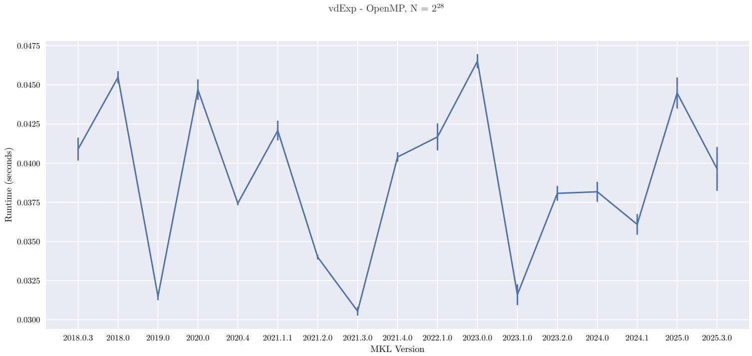

vdExp - Sequential