

# Verifying Hardware Security Modules with Information-Preserving Refinement

Anish Athalye, M. Frans Kaashoek, and Nickolai Zeldovich

Presented by Russell Bentley

Stony Brook


2024



Stony Brook University

# Introduction

- ▶ Presenting [1]
- ▶ Hardware Security Modules (HSM)
- ▶ Theory
- ▶ Implementation
- ▶ Example: TOTP Token



# Hardware Security Modules (HSM)

- ▶ Dedicated hardware for core security functionality
  - ▶ Certificate signing
  - ▶ Password hashing
  - ▶ Token generation
- ▶ May include general purpose core
- ▶ Communicates with host over wire interface



# Thread Model

- ▶ Remote compromise of host
- ▶ Any *digital* use of wire interface
- ▶ Ensure no extra information leaks
  - ▶ Bug in implementation
  - ▶ Timing Attacks →
- ▶ Does not consider physical access
  - ▶ Overvoltage, tampering
  - ▶ Secure as the specification

```
// return error if PIN guess limit exceeded
// ...

// check PIN guess and update guess_count accordingly
if (!constant_time_cmp(&entry->pin, guess)) {
    entry->bad_guesses++;
    uart_write(ERR_BAD_PIN);
    return;
}
entry->bad_guesses = 0;

// output secret
// ...
```




# Information Preserving Refinement (IPR)

- ▶ Start with functional specification
- ▶ Device should implement this spec
- ▶ Wire-level interactions  $\longleftrightarrow$  spec. level operations
- ▶ Based on dual views
  - ▶ Real world (uncompromised host)
  - ▶ Ideal world (compromised host)

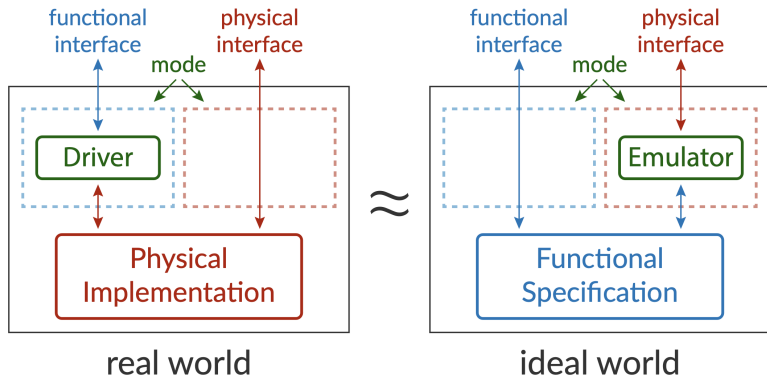
```
var bad_guesses = 0, secret = 0, pin = 0

def store(new_secret, new_pin):
    secret = new_secret
    pin = new_pin
    bad_guesses = 0

def retrieve(guess):
    if bad_guesses >= 10:
        return 'No more guesses'
    if guess == pin:
        bad_guesses = 0
        return secret
    bad_guesses = bad_guesses + 1
    return 'Incorrect PIN'
```



# Information Preserving Refinement (IPR)



# Proving IPR

- ▶ Developer provides  $R$ , relates
  - ▶ State of spec state machine
  - ▶ State of implementation
- ▶ Developer also provides
  - ▶ Driver
  - ▶ Emulator




Figure 7: Functional equivalence: for all implementation states  $c_1$  and spec states  $f_1$  that are related by  $R$ , and for all spec-level operations  $op$ :

- (1) the spec-level output  $v$  matches the driver output
- (2) the final states  $c_2$  and  $f_2$  are related by  $R$




Figure 8: Physical equivalence: for all spec states  $f_1$  and implementation states  $c_1$  that are related by  $R$ , and for all wire-level inputs  $i_1 \dots i_n$ :

- (1) the circuit outputs  $o_1 \dots o_n$  match the emulator outputs
- (2) the final states  $f_2$  and  $c_3$  ( $c_2$  after a reset) are related by  $R$

# Implementation: Knox

- ▶ Knox, monolithic verification
- ▶ Targets Rosette [3]
  - ▶ Solver-aided programming language
- ▶ Generate SMT Constraints
  - ▶ Solve with Z3 [2]



 The Rosette Language




Figure 11: An overview of the Knox workflow. Trusted inputs are shown in green.





# Implementation pt. 2

TODO



# Case Studies

TODO



# Case Study: PIN-protected backup HSM

TODO



# Case Study: Password Hashing HSM

TODO



# Case Study: TOTP Token

TODO



# Discussion

- ▶ Emulator Efficiency
- ▶ Randomness
- ▶ Allowed leakage
- ▶ Monolithic Verification



# Conclusion

TODO



# References I

- [1] Anish Athalye, M Frans Kaashoek, and Nikolai Zeldovich. “Verifying hardware security modules with {Information-Preserving} refinement”. In: *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*. 2022, pp. 503–519.
- [2] Leonardo De Moura and Nikolaj Bjørner. “Z3: An efficient SMT solver”. In: *International conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer. 2008, pp. 337–340.
- [3] Emina Torlak and Rastislav Bodik. “Growing solver-aided languages with rosette”. In: *Proceedings of the 2013 ACM international symposium on New ideas, new paradigms, and reflections on programming & software*. 2013, pp. 135–152.

