

# Verifying Hardware Security Modules with Information-Preserving Refinement

Anish Athalye, M. Frans Kaashoek, and Nickolai Zeldovich

Presented by Russell Bentley

Stony Brook

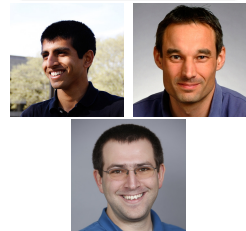
2024



Stony Brook University

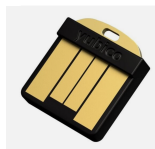
# Introduction

- ▶ Presenting [1] (2022)
- ▶ Hardware Security Modules (HSM)
- ▶ Information Preserving Refinement
- ▶ Knox Framework
- ▶ Case studies



# Hardware Security Modules (HSM)

- ▶ Dedicated hardware for core security functionality
  - ▶ Certificate signing
  - ▶ Password hashing
  - ▶ Token generation
- ▶ High Value Targets
- ▶ Hardware + Software
  - ▶ May include general purpose core
  - ▶ Communicates with host over wire interface



# Threat Model

- ▶ Remote compromise of host
- ▶ *digital* (mis)use of interface
- ▶ Ensure no extra information leaks
  - ▶ Bug in implementation
  - ▶ Timing Attacks →
  - ▶ Crash Safety
- ▶ Does not consider physical access
  - ▶ Overvoltage, tampering
  - ▶ Secure as the specification

```
// return error if PIN guess limit exceeded
// ...

// check PIN guess and update guess_count accordingly
if (!constant_time_cmp(&entry->pin, guess)) {
    entry->bad_guesses++;
    uart_write(ERR_BAD_PIN);
    return;
}
entry->bad_guesses = 0;

// output secret
// ...
```



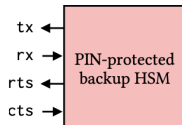
# Information Preserving Refinement (IPR)

- ▶ Start with functional specification
- ▶ Device implements this spec
- ▶ Wire-level interactions  $\longleftrightarrow$  spec. level operations
- ▶ Implementation States  $\longleftrightarrow$  spec. level states
- ▶ Based on dual view
  - ▶ Real world (uncompromised host)
  - ▶ Ideal world (compromised host)

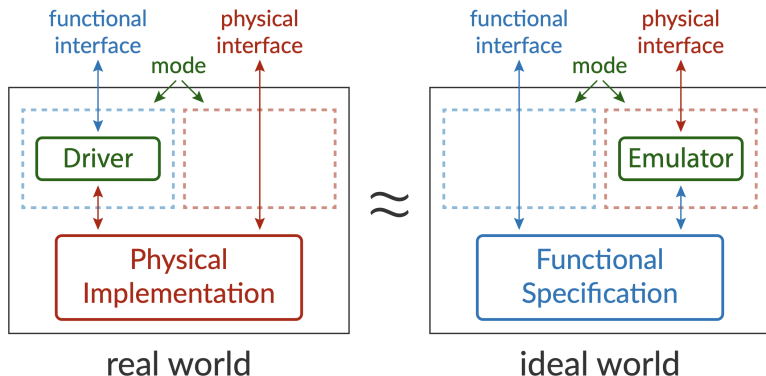
```
var bad_guesses = 0, secret = 0, pin = 0

def store(new_secret, new_pin):
    secret = new_secret
    pin = new_pin
    bad_guesses = 0

def retrieve(guess):
    if bad_guesses >= 10:
        return 'No more guesses'
    if guess == pin:
        bad_guesses = 0
        return secret
    bad_guesses = bad_guesses + 1
    return 'Incorrect PIN'
```



# Information Preserving Refinement (IPR)



# Proving IPR

- ▶ Developer provides  $R$ , relates
  - ▶ State of spec state machine
  - ▶ State of implementation
- ▶ Developer also provides
  - ▶ Driver
  - ▶ Emulator

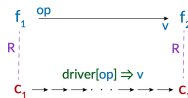


Figure 7: Functional equivalence: for all implementation states  $c_1$  and spec states  $f_1$  that are related by  $R$ , and for all spec-level operations  $op$ :

- (1) the spec-level output  $v$  matches the driver output
- (2) the final states  $c_2$  and  $f_2$  are related by  $R$

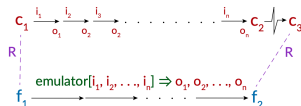


Figure 8: Physical equivalence: for all spec states  $f_1$  and implementation states  $c_1$  that are related by  $R$ , and for all wire-level inputs  $i_1 \dots i_n$ :

- (1) the circuit outputs  $o_1 \dots o_n$  match the emulator outputs
- (2) the final states  $f_2$  and  $c_3$  ( $c_2$  after a reset) are related by  $R$

# Implementation: Knox

- ▶ Knox, monolithic verification
  - ▶ Source on [Github](#)
  - ▶ Hybrid Symbolic Execution [3]
- ▶ Targets Rosette [4]
  - ▶ Solver-aided programming language
- ▶ Generate SMT Constraints
  - ▶ Solve with Z3 [2]



 The Rosette Language

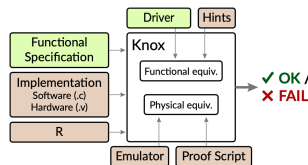


Figure 11: An overview of the Knox workflow. Trusted inputs are shown in green.



Stony Brook University



# Symbolic Execution

- ▶ Functional Equivalence - yield
  - ▶ How to handle Nondeterminism?
  - ▶ Find closure of circuit states
  - ▶ Fork on elements of set
  - ▶ Exponential growth!
  - ▶ *merge* hints shrink space
- ▶ Physical Equivalence - input size
  - ▶ Arbitrary sequence of wire inputs
  - ▶ Intractable to write circuit invariant
  - ▶ *Guided symbolic model checking*
- ▶ Knox allows variety of hints

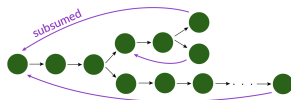


Figure 10: An illustration of *guided symbolic model checking* exploring a state space. Each green circle is a symbolic term representing a set of states. Black arrows show STEP invocations and purple arrows show SUBSUMED invocations.



# Case Studies

- ▶ Three Case Studies
  - ▶ PIN-protected backup
  - ▶ Password hashing
  - ▶ TOTP Token
- ▶ Source on [Github](#)

HSM	Spec		Driver	HW	SW	Proof
	core	total				
PIN backup	32	60	110	2670	190	470
PW hasher	5	150	90	3020	240	650
TOTP	10	180	80	2950	360	830

Figure 16: Lines of code for case studies. Lines of code for the spec are broken down into “core” and “total”, where core is the main HSM functionality and doesn’t include boilerplate or definitions of functions like SHA1, HMAC, and TOTP.

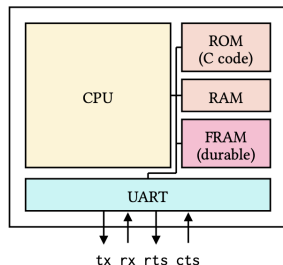
HSM	FE-N	FE-N+C	FE	FE+C	PE
PIN backup	1	10	<b>209</b>	962	<b>8</b>
PW hasher	1	6	<b>74</b>	238	<b>4</b>
TOTP	3	8	<b>44</b>	141	<b>8</b>

Figure 18: Time taken (in minutes) for verification by Knox. FE is functional equivalence; the -N variation disables non-terminism in the driver; +C adds verification of crash safety. PE is physical equivalence. The two bolded columns, FE and PE, together imply IPR.



# Case Study: PIN-protected backup HSM

- ▶ Runs on RISC-V CPU
- ▶ Features
  - ▶ Multiple Secrets
  - ▶ Enforce Attempts Limit
- ▶ Emulator Implementation
  - ▶ Constructive
  - ▶ Inject spec output into circuit
- ▶ Bugs Found
  - ▶ Timing attack with `strcmp`
  - ▶ Guess timing attack



# Case Study: Password Hashing HSM

- ▶ Includes sha256
  - ▶ Needs spec
  - ▶ Verilog implementation
- ▶ Can't retrieve secret
- ▶ Crash Safe
  - ▶ Multiword writes
  - ▶ Caught timing bug
- ▶ Caught mem\_addr bug

```
var secret = 0

def config(new_secret):
    secret = new_secret

def hash(password):
    return sha256(password || secret)
```



# Case Study: TOTP Token

- ▶ Host provides timestamp
- ▶ Can't rewind timestamp
- ▶ Includes SHA1
- ▶ Issue with symbolic execution
- ▶ Timing attack with %

```
var secret = 0, last_timestamp = 0

def set_secret(new_secret):
    secret = new_secret

def get_totp(timestamp):
    if timestamp < last_timestamp:
        return 'Cannot rewind timestamp'
    last_timestamp = timestamp
    return totp(secret, timestamp)

def audit():
    return last_timestamp
```

Figure 14: The functional specification for the TOTP token.

```
/* old implementation:
uint32_t s = (buf[offset] & 0x7f) << 24
| (buf[offset+1] & 0xff) << 16
| (buf[offset+2] & 0xff) << 8
| (buf[offset+3] & 0xff);

*/
uint32_t s = 0;
for (int i = 0; i < 0x10; i++) {
    uint32_t match = ((i != offset) - 1);
    s += ((buf[i] & 0x7f) & match) << 24;
    s += ((buf[i+1] & 0xff) & match) << 16;
    s += ((buf[i+2] & 0xff) & match) << 8;
    s += ((buf[i+3] & 0xff) & match);
}
```

Figure 15: Rewriting TOTP dynamic truncation to avoid symbolic memory addresses.

# Conclusion

- ▶ Today we covered:
  - ▶ HSMs
  - ▶ Information Preserving Refinement
  - ▶ Knox
- ▶ Questions?



# References I

- [1] Anish Athalye, M Frans Kaashoek, and Nickolai Zeldovich. “Verifying hardware security modules with {Information-Preserving} refinement”. In: *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*. 2022, pp. 503–519.
- [2] Leonardo De Moura and Nikolaj Bjørner. “Z3: An efficient SMT solver”. In: *International conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer. 2008, pp. 337–340.
- [3] Emina Torlak and Rastislav Bodik. “A lightweight symbolic virtual machine for solver-aided host languages”. In: *ACM SIGPLAN Notices* 49.6 (2014), pp. 530–541.
- [4] Emina Torlak and Rastislav Bodik. “Growing solver-aided languages with rosette”. In: *Proceedings of the 2013 ACM international symposium on New ideas, new paradigms, and reflections on programming & software*. 2013, pp. 135–152.

